

Object Segmentation to prevent contextual shortcut learning

Robin Szymanski
robin.szymanski.001@student.uni.lu
University of Luxembourg
Luxembourg

Abstract

This project will focus on image classification using a deep learning model trained on a 10-class subset of the iNaturalist-12k dataset available on Kaggle, provided by Aryan Pandey. The dataset contains 12,000 images across biological categories such as Amphibia, Mammalia, and Plantae. The classes are fairly balanced with 1200 available images for each of the classes. The main goal is to design, train, and evaluate a convolutional neural network from scratch on the aforementioned dataset. As a possible extension, the project may explore feature engineering techniques such as extraction of the biological object (animal or plant) to study the impact of background information on classification. Additional extensions include the creation of a second model by transfer learning and the comparison of the latter with the model created from scratch.

1 Introduction

Deep learning models achieve strong performance on standard benchmarks, yet their limitations remain an active area of research. Geirhos et al. define shortcuts as decision rules that perform well on benchmark datasets but fail under more challenging or out-of-distribution conditions [2]. Such shortcut learning is particularly problematic in safety-critical applications such as autonomous driving, where models must generalize to unseen environments [1]. At the same time, prior work has shown that image backgrounds can contain highly predictive information and may substantially contribute to classification accuracy [11].

This project investigates whether reducing background information through automatic object segmentation can improve robustness and reduce shortcut learning in biological image classification. A 10-class subset of the iNaturalist-12K dataset from Kaggle was used [6]. The dataset is well suited for this task because it contains substantial variability in habitats, lighting conditions, and objects. A segmentation pipeline based on Segment Anything (SAM) [5, 8] and CLIP [7] was developed to automatically extract the foreground object without requiring manual annotations. The resulting segmented masks were used to crop images to the region of interest and to train CNN-based classification models. The performance of models trained with this segmentation-based regularization approach was then compared against equivalent models trained on the original, non-segmented images.

2 Related work

Shortcut learning has emerged as an important limitation of deep neural networks [2]. In ecological datasets, background information and hence the location can become a strong shortcut because images from the same location often share very similar environmental context, particularly in datasets collected from fixed animal camera traps [1]. Geirhos et al. argue that shortcut learning should not only

be studied through in-distribution accuracy, but rather by manipulating the evidence available to the model, such as background, texture, or color, and observing how predictions change [2]. Several studies demonstrate that image backgrounds alone can contain sufficient information for classification [11, 12]. Xiao et al. showed that models trained on ImageNet often strongly depend on contextual background information, while Rosenfeld et al. demonstrated that even small background perturbations such as noise can significantly alter predictions [10, 11]. Geirhos et al. further found that CNNs frequently rely more on texture cues than on object shape, which may contribute to such behavior and provides additional motivation for segmentation-based approaches that emphasize object structure over contextual texture information [3].

These findings motivate approaches that explicitly separate foreground objects from their surrounding context. However, many supervised segmentation methods, including segmentation-based variants of YOLO, require manually annotated masks and labels for training, which is highly labor-intensive and not given for the iNaturalist dataset [9]. Such annotations are not available for the iNaturalist-12K dataset used in this project. However, Segment Anything (SAM) [5, 8] provides a way to automatically generate segmentation proposals without requiring manual annotation. To semantically identify relevant masks, this project additionally uses CLIP [7], which uses text-image similarity to semantically score and label images. The resulting pipeline is conceptually related to copy-paste augmentation methods [4], where segmented foreground objects are inserted into alternative backgrounds to improve robustness and reduce contextual bias.

3 Experiments

3.1 Dataset

The dataset used for this project is a 10-class subset of the iNaturalist-12k dataset available on Kaggle (link in appendix), provided by Aryan Pandey. It contains 12,000 images across biological categories such as Amphibia, Mammalia, and Plantae. The classes are balanced with 1,200 available images for each of the classes. The dataset was split into 8,000 training images, 2,000 validation images, and 2000 test images, stratified by class.

3.2 Model Architecture

3.3 Training Details

3.4 Results

The quality of the segmentation masks was manually evaluated on a sample of 200 masks, with 20 masks per class. A mask was counted as good if it captured the main foreground object without substantial missing regions or excessive background. The results in table 1 show significant variation between classes. In particular,

masks were weak for *Animalia*, but very reliable for *Aves*, *Mammalia*, and *Fungi*. Therefore, the evaluation will additionally inspect whether classes with higher-quality masks also benefit more from segmentation in terms of macro-accuracy and AUC.

4 Discussion

- Talking about limitations: Commenting that this is just a student project, that the segmentation is very imperfect, that there were not enough resources to properly test all models. Not enough data either.
- Quality of masks not great
- If results worse with segmentation: Realm of study project quite limited, worth repeating under cleaner conditions with proper labelling. Pipeline quite unreliable.
- Could go much further by putting e.g. a cow in the sky, which is related to what [3] and [10] found.
- Experiment worth to be better repeated for datasets with given segmentation masks. Maybe label them by company.

5 References

References

- [1] Sara Beery, Grant Van Horn, and Pietro Perona. 2018. Recognition in Terra Incognita. In *Proceedings of the European Conference on Computer Vision (ECCV)*.
- [2] Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A Wichmann. 2020. Shortcut learning in deep neural networks. *Nature Machine Intelligence* 2, 11 (2020), 665–673.
- [3] Robert Geirhos, Patricia Rubisch, Claudio Michaelis, Matthias Bethge, Felix A. Wichmann, and Wieland Brendel. 2019. ImageNet-trained CNNs are biased towards texture; increasing shape bias improves accuracy and robustness. In *International Conference on Learning Representations*. <https://openreview.net/forum?id=Bygh9j09KX>
- [4] Golnaz Ghiasi, Yin Cui, Aravind Srinivas, Rui Qian, Tsung-Yi Lin, Ekin D. Cubuk, Quoc V. Le, and Barret Zoph. 2021. Simple Copy-Paste is a Strong Data Augmentation Method for Instance Segmentation. *arXiv:2012.07177 [cs.CV]* <https://arxiv.org/abs/2012.07177>
- [5] Alexander Kirillov, Eric Mintun, Nikhila Ravi, Hanzi Mao, Chloe Rolland, Laura Gustafson, Tete Xiao, Spencer Whitehead, Alexander C. Berg, Wan-Yen Lo, Piotr Dollár, and Ross Girshick. 2023. Segment Anything. *arXiv:2304.02643 (2023)*.
- [6] Aryan Pandey. [n. d.]. iNaturalist-12K Dataset. <https://www.kaggle.com/datasets/aryanpandey1109/inaturalist12k>. Accessed: 2026-05-24.
- [7] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. 2021. Learning Transferable Visual Models From Natural Language Supervision. *CoRR abs/2103.00020 (2021)*. *arXiv:2103.00020* <https://arxiv.org/abs/2103.00020>
- [8] Nikhila Ravi, Valentin Gabeur, Yuan-Ting Hu, Ronghang Hu, Chaitanya Ryali, Tengyu Ma, Haitham Khedr, Roman Rädle, Chloe Rolland, Laura Gustafson, Eric Mintun, Junting Pan, Kalyan Vasudev Alwala, Nicolas Carion, Chao-Yuan Wu, Ross Girshick, Piotr Dollár, and Christoph Feichtenhofer. 2024. SAM 2: Segment Anything in Images and Videos. *arXiv:2408.00714 [cs.CV]* <https://arxiv.org/abs/2408.00714>
- [9] Joseph Redmon, Santosh Kumar Divvala, Ross B. Girshick, and Ali Farhadi. 2015. You Only Look Once: Unified, Real-Time Object Detection. *CoRR abs/1506.02640 (2015)*. *arXiv:1506.02640* <http://arxiv.org/abs/1506.02640>
- [10] Amir Rosenfeld, Richard Zemel, and John K. Tsotsos. 2018. The Elephant in the Room. *arXiv:1808.03305 [cs.CV]* <https://arxiv.org/abs/1808.03305>
- [11] Kai Yuanqing Xiao, Logan Engstrom, Andrew Ilyas, and Aleksander Madry. 2021. Noise or Signal: The Role of Image Backgrounds in Object Recognition. In *International Conference on Learning Representations*. <https://openreview.net/forum?id=gl3D-xY7wLq>
- [12] Zhuotun Zhu, Lingxi Xie, and Alan L. Yuille. 2016. Object Recognition with and without Objects. *CoRR abs/1611.06596 (2016)*. *arXiv:1611.06596* <http://arxiv.org/abs/1611.06596>

A Dataset

<https://www.kaggle.com/datasets/aryanpandey1109/inaturalist12k>

Table 1: Manual quality assessment of segmentation masks by class (n=200). Values indicate the percentage of masks rated as good.

Class	Amph.	Anim.	Arach.	Aves	Fungi
Good masks (%)	50	29	43	94	79
Class	Insecta	Mamm.	Moll.	Plant.	Rept.
Good masks (%)	74	81	74	76	67

B Results

C Declaration of generative AI usage

I hereby declare that I have used an LLM (ChatGPT, model GPT-5.3 and GPT-5.5) to assist me with some small passages of the code, debugging, improve some readability, to get new ideas, to assist me with training on the HPC server, and to find better vocabulary and restructuring some sentences while writing the paper. This is in accordance with the "Guidelines on the Use of Generative AI for Teaching and Learning", Version: 1.0, Date: 2026-02-16.